

Introduction to Deep Learning (CS474)

Lecture 2

Outline

- Pytorch Basics.
 - **Tensor**
 - Example
 - Numeric type
 - Operations
 - NumPy bridge

PyTorch Basics.

- PyTorch gives us a data type, the **Tensor**, to hold numbers, vectors, matrices, or arrays in general.
- In addition, it provides functions for operating on them.
- We can program with them incrementally and, if we want, interactively, just like we are used to from Python. If you know NumPy, this will be very familiar.

Tensor




Figure 1: Tensors are the building blocks for representing data in PyTorch

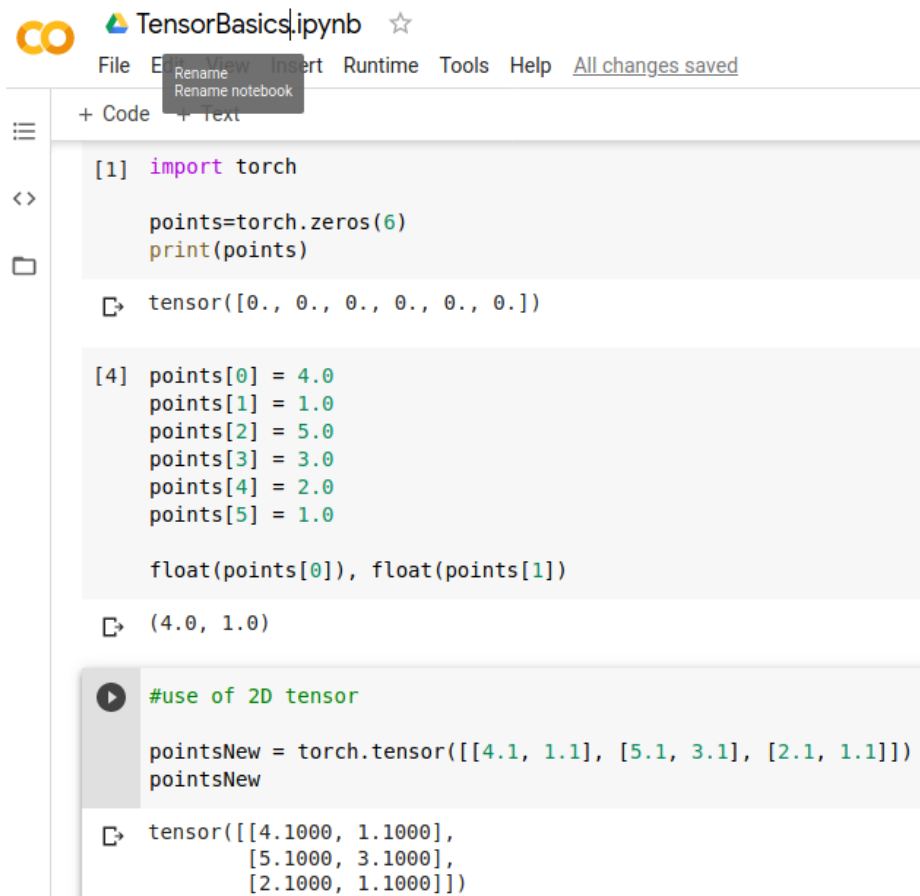
Tensor

- PyTorch is not the only library that deals with multidimensional arrays.
- Compared to NumPy arrays, PyTorch tensors have a few superpowers.
- PyTorch tensors have the ability to perform very fast operations on graphical processing units (GPUs), distribute operations on multiple devices or machines, and keep track of the graph of computations that created them.
- These are all important features when implementing a modern deep learning library.

Tensor: Example

- We would like to use to represent a geometrical object.
- Let us consider a 2D triangle with vertices at coordinates (4, 1), (5, 3), and (2, 1).
- The example is not particularly pertinent to deep learning, but it's easy to follow.
- Instead of having coordinates as numbers in a Python list, as we did earlier, we can use a one-dimensional tensor by storing Xs in the even indices and Ys in the odd indices.

Tensor: Example (contd.)



The screenshot shows a Jupyter Notebook interface with the title 'TensorBasics.ipynb'. The notebook has three code cells. The first cell imports 'torch' and creates a 1D tensor of zeros. The second cell updates the 'points' list with specific values and prints the first two elements. The third cell creates a 2D tensor from a list of lists.

```
[1] import torch

points=torch.zeros(6)
print(points)

tensor([0., 0., 0., 0., 0., 0.])

[4] points[0] = 4.0
points[1] = 1.0
points[2] = 5.0
points[3] = 3.0
points[4] = 2.0
points[5] = 1.0

float(points[0]), float(points[1])

(4.0, 1.0)

#use of 2D tensor

pointsNew = torch.tensor([[4.1, 1.1], [5.1, 3.1], [2.1, 1.1]])
pointsNew

tensor([[4.1000, 1.1000],
        [5.1000, 3.1000],
        [2.1000, 1.1000]])
```

Tensor: Example (contd.)

```
[10] # We can ask the tensor about its shape!  
     # This informs us about the size of the tensor along each dimension
```

```
pointsNew.shape
```

```
↳ torch.Size([3, 2])
```

```
[12] # Now we can access an individual element in the tensor using two indices:  
     pointsNew[0, 1]  
     #Above line help us to get the Y-coordinate of the zeroth point
```

```
↳ tensor(1.1000)
```

```
▶ # We can also access the first element in the tensor.  
  pointsNew[0]
```

```
↳ tensor([4.1000, 1.1000])
```


Tensor: Numeric types

- So far, we have covered the basics of how tensors work, but we have not yet touched on **what kinds of numeric types** we can **store in a Tensor**.
- The **dtype** argument to tensor constructors (that is, functions like `tensor` , `zeros` , and `ones`) specifies the numerical data (d) type that will be contained in the tensor.
- possible values for the dtype argument:
 - `torch.float32` or `torch.float`: 32-bit floating-point
 - `torch.float64` or `torch.double`: 64-bit, double-precision floating-point
 - `torch.float16` or `torch.half`: 16-bit, half-precision floating-point
 - `torch.int8`: signed 8-bit integers
 - `torch.uint8`: unsigned 8-bit integers
 - `torch.int16` or `torch.short`: signed 16-bit integers
 - `torch.int32` or `torch.int`: signed 32-bit integers
 - `torch.int64` or `torch.long`: signed 64-bit integers
 - `torch.bool`: Boolean

Tensor: Numeric types

- Computations happening in neural networks are typically executed with 32-bit floating-point precision.
- Higher precision, like 64-bit, will not buy improvements in the accuracy of a model and will require more memory and computing time.

```
▶ import torch

#In order to allocate a tensor of the right numeric type,
#we can specify the proper dtype as an argument to the constructor.

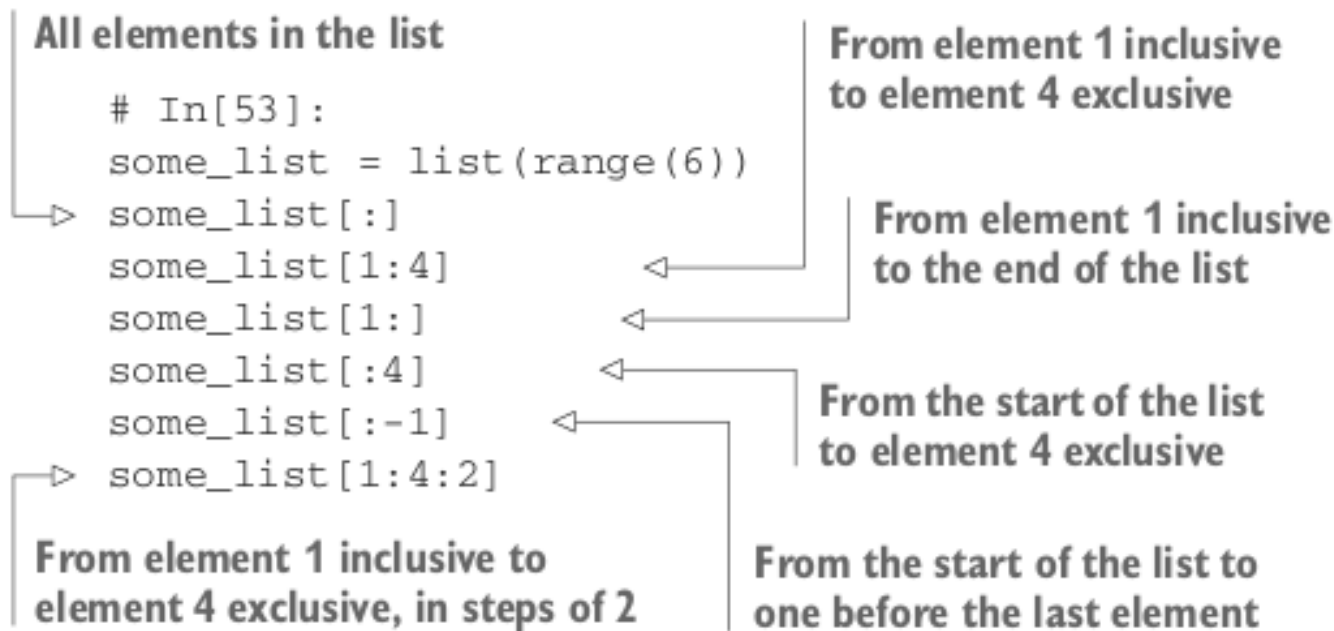
double_points=torch.ones(10, 2, dtype=torch.double)
short_points=torch.tensor([[1, 2], [22, 33]], dtype=torch.short)

print(double_points)
print(short_points)
```

```
↳ tensor([[1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.],
          [1., 1.]], dtype=torch.float64)
tensor([[ 1,  2],
        [22, 33]], dtype=torch.int16)
```

Tensor: Operations (Indexing tensors)

- Reminder!



Tensor: Operations (Indexing tensors)

- Hope you remember our previous *Notebook* example with **pointsNew**!

```
[14] # All rows after the first; implicitly all columns.
```

```
pointsNew[1:]
```

```
↳ tensor([[5.1000, 3.1000],  
          [2.1000, 1.1000]])
```



```
#All rows after the first; all columns
```

```
pointsNew[1:, :]
```

```
↳ tensor([[5.1000, 3.1000],  
          [2.1000, 1.1000]])
```

Tensor: Operations (transpose)

- You can create a *Notebook* in Google colab for experimenting with the tensor API.
- Vast majority of operations on and between tensors are available in the **torch module**.

```
# Checking transpose.  
  
a=torch.ones(3, 2)  
a_transpose=torch.transpose(a, 0, 1)  
  
print(a.shape)  
print(a_transpose.shape)  
  
torch.Size([3, 2])  
torch.Size([2, 3])
```

Tensor: Operations (addition)

```
▶ import torch

# Initialize

x = torch.ones(2, 3)
y = torch.rand(2, 3)    # Initialize with random values

# Operations

z1 = x + y
z2 = torch.add(x, y)

print(z2)
print(z1)
```

```
↳ tensor([[1.2826, 1.3097, 1.0306],
          [1.6098, 1.7936, 1.0860]])
tensor([[1.2826, 1.3097, 1.0306],
          [1.6098, 1.7936, 1.0860]])
```

Tensor: Operations (In-place operations)

- Certain operations exist only as methods of the Tensor object.
- They are recognizable from a trailing underscore in their name, like `zero_`, which indicates that the method operates in place by modifying the input instead of creating a new output tensor and returning it.
- For instance, the `zero_` method zeros out all the elements of the input.

```
# checking In-place operation.  
new = torch.ones(3, 2)  
new.zero_()  
print(new)
```

```
tensor([[0., 0.],  
        [0., 0.],  
        [0., 0.]])
```

Tensor: NumPy bridge

- Converting a Torch Tensor to a NumPy array and vice versa is a breeze.
 - The Torch Tensor and NumPy array will share their underlying memory locations (if the Torch Tensor is on CPU), and changing one will change the other.
-
- Converting a Torch Tensor to a NumPy Array
 - Converting NumPy Array to Torch Tensor

Tensor: NumPy bridge

- Converting a Torch Tensor to a NumPy Array

```
[2] import torch
import numpy
# checking bridge!
a = torch.ones(5)
print(a)
b = a.numpy()
print(b)
```

```
↳ tensor([1., 1., 1., 1., 1.])
[1. 1. 1. 1. 1.]
```

```
▶ #Now we will see how the numpy array changed in value.
a.add_(1)
print(a)
print(b)
```

```
↳ tensor([2., 2., 2., 2., 2.])
[2. 2. 2. 2. 2.]
```

Tensor: NumPy bridge

- Converting NumPy Array to Torch Tensor

```
import numpy as np
a = np.ones(5)
b = torch.from_numpy(a)
np.add(a, 1, out=a)
print(a)
print(b)
```

Out:

```
[2. 2. 2. 2. 2.]
tensor([2., 2., 2., 2., 2.], dtype=torch.float64)
```

References

- All the contents present in the slides are taken from various online resources. Due credit is given in the respective slides. These slides are used for *academic* purposes only.