

Introduction to Deep Learning (CS474)

Lecture 6

Outline

- Mechanics of Learning-PART I
 - Introduction
 - Example
 - Application of Linear Model.

Introduction

- We can argue that **learning from data** presumes the underlying model is not engineered to solve a specific problem and is instead capable of **approximating** a much wider family of functions.
- We're interested in models that are not engineered for solving a specific narrow task, but that can be automatically adapted to specialize themselves for any one of many similar tasks using input and output pairs.
- **PyTorch** is designed to make it easy to create models for which the derivatives of the fitting error, with respect to the parameters, can be expressed analytically.

Introduction: Overview of Learning Process

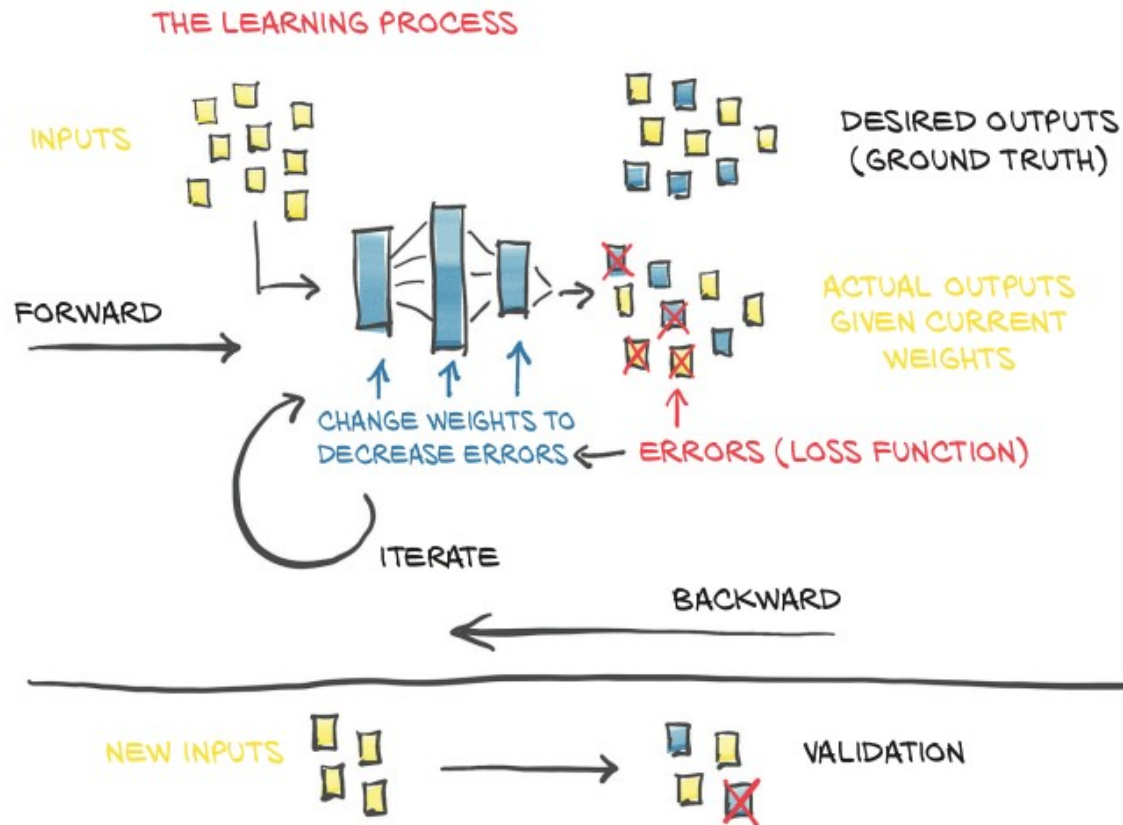


Image credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

Example

```
▶ import numpy as np
import torch

# t_c values are temperatures in Celsius
# t_u values are our unknown units

t_c = [0.5, 14.0, 15.0, 28.0, 11.0, 8.0, 3.0, -4.0, 6.0, 13.0, 21.0]
t_u = [35.7, 55.9, 58.2, 81.9, 56.3, 48.9, 33.9, 21.8, 48.4, 60.4, 68.4]

# For convenience, we've already put the data into tensors.

t_c = torch.tensor(t_c)
t_u = torch.tensor(t_u)
```

Example: Application of Linear Model

- The **two** may be **linearly** related—that is, multiplying t_u by a factor and adding a constant, we may get the temperature in Celsius (up to an error that we omit):

$$t_c = w * t_u + b$$

- We'll see how well the final model performs.
- We chose to name **w** and **b** after **weight** and **bias**, two very common terms for *linear scaling* and the *additive constant*.

Example: Application of Linear Model

- We have a **model** with some unknown parameters.
- We need to **estimate** those parameters so that the error between predicted outputs and measured values is as low as possible.
- We notice that we still need to exactly define a measure of the error. Such a measure, which we refer to as the loss function, should be high if the error is high and should ideally be as low as possible for a perfect match.
- Our **optimization process** should therefore aim at finding \mathbf{w} and \mathbf{b} so that the **loss function** is at a minimum.

Example: Application of Linear Model

- Continuing with the earlier *notebook*.

```
# Our model!
def model(t_u, w, b):
    return w * t_u + b

# Our loss function!
def loss_fn(t_p, t_c):
    squared_diffs = (t_p - t_c)**2
    return squared_diffs.mean()
```


Example: Application of Linear Model

- Note that we are building a **tensor** of differences, taking their square **element-wise**, and finally producing a scalar loss function by averaging all of the elements in the resulting tensor. It is a mean square loss.
- Continuing with the earlier *notebook*.

```
# Initializing parameters!
```

```
w = torch.ones(())
```

```
b = torch.zeros(())
```

```
t_p = model(t_u, w, b)
```

```
t_p
```

```
↳ tensor([35.7000, 55.9000, 58.2000, 81.9000, 56.3000, 48.9000, 33.9000, 21.8000,  
         48.4000, 60.4000, 68.4000])
```

Example: Application of Linear Model

- Now, we can print the loss as well!
- Continuing with the earlier *notebook*.

```
# Loss calculation!
```

```
loss = loss_fn(t_p, t_c)  
loss
```

```
↳ tensor(1763.8846)
```

- How do we estimate \mathbf{w} and \mathbf{b} such that the loss reaches a **minimum**?

Example: Application of Linear Model



Figure 1 A cartoon depiction of the optimization process, where a person with knobs for w and b searches for the direction to turn the knobs that makes the loss decrease

Image credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

Example: Application of Linear Model

Decreasing Loss:

- **Gradient descent** is not that different from the scenario we just described.
- The idea is to compute the rate of change of the loss with respect to each parameter, and modify each parameter in the direction of decreasing loss.

```
delta = 0.1

loss_rate_of_change_w = \
    (loss_fn(model(t_u, w + delta, b), t_c) -
     loss_fn(model(t_u, w - delta, b), t_c)) / (2.0 * delta)
```

Example: Application of Linear Model

- We typically should scale the rate of change by a small factor.
- This scaling factor has many names; the one we use in machine learning is `learning_rate`.

```
learning_rate = 1e-2  
  
print(learning_rate)  
  
w = w - learning_rate * loss_rate_of_change_w
```

```
0.01
```

Example: Application of Linear Model

- This represents the basic **parameter-update step** for **gradient descent**.

```
# Same thing can be done for b

loss_rate_of_change_b = \
    (loss_fn(model(t_u, w, b + delta), t_c) -
     loss_fn(model(t_u, w, b - delta), t_c)) / (2.0 * delta)

b = b - learning_rate * loss_rate_of_change_b
```

Example: Application of Linear Model

```
#Computing the derivative

def dloss_fn(t_p, t_c):
    dsq_diffs = 2 * (t_p - t_c) / t_p.size(0)
    return dsq_diffs

# APPLYING THE DERIVATIVES TO THE MODEL

def dmodel_dw(t_u, w, b):
    return t_u

def dmodel_db(t_u, w, b):
    return 1.0

# DEFINING THE GRADIENT FUNCTION

def grad_fn(t_u, t_c, t_p, w, b):
    dloss_dtp = dloss_fn(t_p, t_c)
    dloss_dw = dloss_dtp * dmodel_dw(t_u, w, b)
    dloss_db = dloss_dtp * dmodel_db(t_u, w, b)
    return torch.stack([dloss_dw.sum(), dloss_db.sum()])
```

Slide credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

Example: Application of Linear Model

- We now have everything in place to optimize our parameters.
- Starting from a tentative value for a parameter, we can **iteratively** apply updates to it for a fixed number of iterations, or until w and b stop changing.
- There are several stopping criteria; for now, we'll stick to a fixed number of iterations.
- We call a training iteration during which we update the parameters for all of our training samples an **epoch**.

Example: Application of Linear Model

```
# Training Loop

def training_loop(n_epochs, learning_rate, params, t_u, t_c):
    for epoch in range(1, n_epochs + 1):
        w, b = params

        t_p = model(t_u, w, b) # <1>
        loss = loss_fn(t_p, t_c)
        grad = grad_fn(t_u, t_c, t_p, w, b) # <2>

        params = params - learning_rate * grad

        print('Epoch %d, Loss %f' % (epoch, float(loss))) # <3>

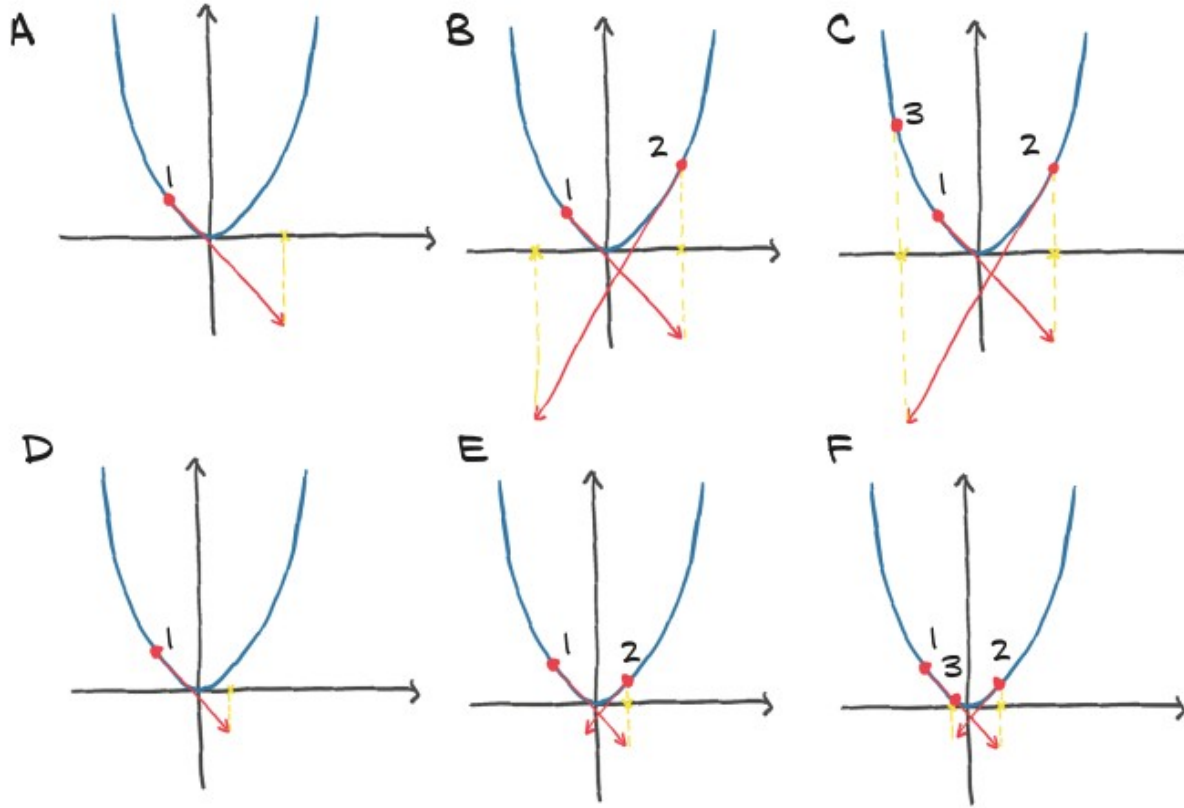
    return params
```

Example: Application of Linear Model

```
# Invoking Training Loop

training_loop(
    n_epochs = 100,
    learning_rate = 1e-2,
    params = torch.tensor([1.0, 0.0]),
    t_u = t_u,
    t_c = t_c)
```

Example: Application of Linear Model



Slide credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

References

- All the contents present in the slides are taken from various online resources. Due credit is given in the respective slides. These slides are used for *academic* purposes only.