

Introduction to Deep Learning (CS474)

Lecture 7

Outline

- Mechanics of Learning-PART II
 - Pytorch's autograd

Introduction

- In our little adventure, we just saw a simple example of backpropagation: we computed the gradient of a composition of functions—the model and the loss—with respect to their innermost parameters (\mathbf{w} and \mathbf{b}) by propagating derivatives backward using the chain rule.
- The basic requirement here is that **all functions** we're dealing with can be **differentiated** analytically.
- Writing the analytical expression for the derivatives of a very deep composition of linear and nonlinear functions is not a lot of fun.

Introduction

- This is when PyTorch tensors come to the rescue, with a PyTorch component called **autograd**.
- We left out one very interesting aspect, however: PyTorch tensors can remember where they come from, in terms of the operations and parent tensors that originated them, and they can automatically provide the chain of derivatives of such operations with respect to their inputs.
- This means we won't need to derive our model by **hand**; given a forward expression, no matter how nested, PyTorch will automatically provide the gradient of that expression with respect to its input parameters.

Example



```
import numpy as np
import torch

t_c = torch.tensor([0.5, 14.0, 15.0, 28.0, 11.0, 8.0,
                    3.0, -4.0, 6.0, 13.0, 21.0])
t_u = torch.tensor([35.7, 55.9, 58.2, 81.9, 56.3, 48.9,
                    33.9, 21.8, 48.4, 60.4, 68.4])

t_un = 0.1 * t_u

def model(t_u, w, b):
    return w * t_u + b

def loss_fn(t_p, t_c):
    squared_diffs = (t_p - t_c)**2
    return squared_diffs.mean()
```

Example: Applying AUTOGRAD

Continuing with the earlier notebook:

```
params = torch.tensor([1.0, 0.0], requires_grad=True)
```

- Notice the `requires_grad=True` argument to the tensor constructor? That argument is telling PyTorch to track the entire family tree of tensors resulting from operations on `params` .
- In other words, any tensor that will have `params` as an ancestor will have access to the chain of functions that were called to get from `params` to that tensor.
- In case these functions are differentiable (and most PyTorch tensor operations will be), the value of the derivative will be automatically populated as a `grad` attribute of the `params` tensor.
- In general, all PyTorch tensors have an attribute named `grad` .

Example: Applying AUTOGRAAD

Continuing with the earlier notebook:

```
loss = loss_fn(model(t_u, *params), t_c)
loss.backward()
```

```
params.grad
```

```
↳ tensor([4517.2969,  82.6000])
```

- We have called the *model* and compute the *loss*, and then call **backward** on the loss tensor.

Example: Applying AUTOGRAD

- When we compute our loss while the parameters w and b require gradients, in addition to performing the actual computation, PyTorch creates the **autograd graph** with the operations (in black circles) as nodes, as shown in the figure 1.

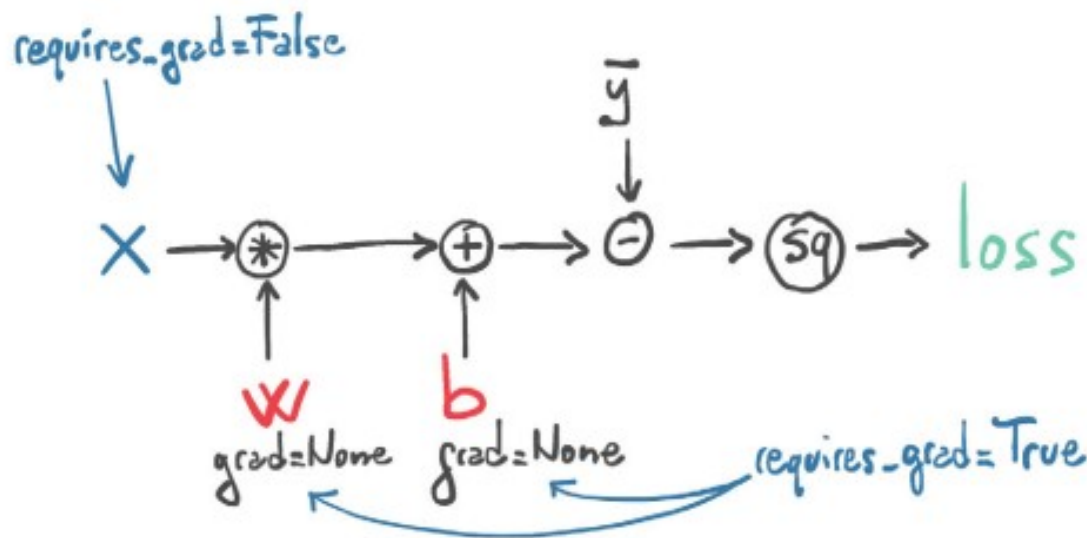


Figure 1

Slide & Image credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

Example: Applying AUTOGRAD

- When we call `loss.backward()`, PyTorch traverses this graph in the **reverse** direction to compute the gradients, as shown by the arrows in the figure 2.




Figure 2

Example: Applying AUTOGRAD

- Calling backward will lead derivatives to accumulate at leaf nodes. We need to zero the gradient explicitly after using it for parameter updates.
- We can do this easily using the in-place `zero_` method.
- Therefore, continuing with the earlier *notebook*.

```
if params.grad is not None:  
    params.grad.zero_()
```

Example: Applying AUTOGRAD

```
# Our autograd-enabled training code

def training_loop(n_epochs, learning_rate, params, t_u, t_c):
    for epoch in range(1, n_epochs + 1):
        if params.grad is not None: # <1>
            params.grad.zero_()

        t_p = model(t_u, *params)
        loss = loss_fn(t_p, t_c)
        loss.backward()

        with torch.no_grad(): # <2>
            params -= learning_rate * params.grad

        if epoch % 500 == 0:
            print('Epoch %d, Loss %f' % (epoch, float(loss)))

    return params
```

Slide credit: E. STEVENS, L. ANTIGA, and T. VIEHMANN

Example: Applying AUTOGRAAD



Invoking training loop

```
training_loop(  
    n_epochs = 5000,  
    learning_rate = 1e-2,  
    params = torch.tensor([1.0, 0.0], requires_grad=True),  
    t_u = t_un,  
    t_c = t_c)
```



```
Epoch 500, Loss 7.860116  
Epoch 1000, Loss 3.828538  
Epoch 1500, Loss 3.092191  
Epoch 2000, Loss 2.957697  
Epoch 2500, Loss 2.933134  
Epoch 3000, Loss 2.928648  
Epoch 3500, Loss 2.927830  
Epoch 4000, Loss 2.927679  
Epoch 4500, Loss 2.927652  
Epoch 5000, Loss 2.927647  
tensor([ 5.3671, -17.3012], requires_grad=True)
```

References

- All the contents present in the slides are taken from various online resources. Due credit is given in the respective slides. These slides are used for *academic* purposes only.